

Assignment 8

#

In this assignment you'll explore the relationship between model complexity and generalization performance, by adjusting key parameters of various supervised learning models.

Classification

#

Here's an application of machine learning that could save your life! For this section of the assignment we will be working with the [UCI Mushroom Data Set](stored in `mushrooms.csv`. The data will be used to train a model to predict whether or not a mushroom is poisonous.

The following attributes are provided:

#

Attribute Information:

#

1. cap-shape: bell=b, conical=c, convex=x, flat=f, knobbed=k, sunken=s

2. cap-surface: fibrous=f, grooves=g, scaly=y, smooth=s

3. cap-color: brown=n, buff=b, cinnamon=c, gray=g, green=r, pink=p, purple=u, red=e, white=w, yellow=y

4. bruises?: bruises=t, no=f

5. odor: almond=a, anise=l, creosote=c, fishy=y, foul=f, musty=m, none=n, pungent=p, spicy=s

6. gill-attachment: attached=a, descending=d, free=f, notched=n

7. gill-spacing: close=c, crowded=w, distant=d

8. gill-size: broad=b, narrow=n
9. gill-color: black=k, brown=n, buff=b, chocolate=h,
gray=g, green=r, orange=o, pink=p, purple=u, red=e,
white=w, yellow=y
10. stalk-shape: enlarging=e, tapering=t
11. stalk-root: bulbous=b, club=c, cup=u, equal=e,
rhizomorphs=z, rooted=r, missing=?
12. stalk-surface-above-ring: fibrous=f, scaly=y,
silky=k, smooth=s
13. stalk-surface-below-ring: fibrous=f, scaly=y,
silky=k, smooth=s
14. stalk-color-above-ring: brown=n, buff=b,
cinnamon=c, gray=g, orange=o, pink=p, red=e, white=w,
yellow=y
15. stalk-color-below-ring: brown=n, buff=b,
cinnamon=c, gray=g, orange=o, pink=p, red=e, white=w,
yellow=y
16. veil-type: partial=p, universal=u
17. veil-color: brown=n, orange=o, white=w, yellow=y
18. ring-number: none=n, one=o, two=t
19. ring-type: cobwebby=c, evanescent=e, flaring=f,
large=l, none=n, pendant=p, sheathing=s, zone=z
20. spore-print-color: black=k, brown=n, buff=b,
chocolate=h, green=r, orange=o, purple=u, white=w,
yellow=y
21. population: abundant=a, clustered=c, numerous=n,
scattered=s, several=v, solitary=y
22. habitat: grasses=g, leaves=l, meadows=m, paths=p,
urban=u, waste=w, woods=d

#

- The data in the mushrooms dataset is currently encoded with strings.
- These values will need to be encoded to numeric to work with sklearn.
- We'll use `pd.get_dummies` to convert the categorical variables into indicator variables.

```
# In[4]:
```

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

mush_df = pd.read_csv('mushrooms.csv')
mush_df2 = pd.get_dummies(mush_df)

X_mush = mush_df2.iloc[:,2:]
y_mush = mush_df2.iloc[:,1]

# use the variables X_train2, y_train2 for Question 1
X_train2, X_test2, y_train2, y_test2 =
train_test_split(X_mush, y_mush, random_state=0)

# For performance reasons in Questions 2 , we will
create a smaller version of the
# entire mushroom dataset for use in those questions.
For simplicity we'll just re-use
# the 25% test split created above as the
representative subset.
#
# Use the variables X_subset, y_subset for Questions 2
X_subset = X_test2
y_subset = y_test2
```

```
# ### Question 1
```

```
#
```

```
# Using `X_train2` and `y_train2` from the preceeding  
cell, train a DecisionTreeClassifier with default  
parameters and random_state=0.
```

What are the 5 most important features found by the
decision tree?

```
#
```

```
# As a reminder, the feature names are available in the  
`X_train2.columns` property, and the order of the  
features in `X_train2.columns` matches the order of the  
feature importance values in the classifier's  
`feature_importances_` property.
```

```
#
```

```
# *This function should return a list of length 5  
containing the feature names in descending order of  
importance.*
```

```
#
```

```
# In[102]:
```

```
def answer_one():
```

```
    from sklearn.tree import DecisionTreeClassifier
```

```
    tree_clf = DecisionTreeClassifier().fit(X_train2,  
y_train2)
```

```
    feature_names = []
```

```
    # Get index of importance leaves since theirs order  
is the same with feature columns
```

```
    Write ur code here
```

```
        # Add importance so we can further order this
list, and add feature name with index
```

```
    Write ur code here
```

```
    # Descending sort
```

```
    Write ur code here
```

```
    # Turn in to a numpy array
```

```
Write ur code here
```

```
    # Select only feature names
```

```
    Write ur code here
```

```
    # Turn back to python list
```

```
    Write ur code here
```

```
    return feature_names # Your answer here
```

```
answer_one()
```

```
# ### Question 2
```

```
#  
# For this question, we're going to use the  
`validation_curve` function in  
`sklearn.model_selection` to determine training and  
test scores for a Support Vector Classifier (`SVC`)  
with varying parameter values. In the validation_curve  
function, in addition to taking an initialized unfitted  
classifier object, takes a dataset as input and does  
its own internal train-test splits to compute results.  
#  
# **Because creating a validation curve requires  
fitting multiple models, for performance reasons this  
question will use just a subset of the original  
mushroom dataset:
```

```
please use the variables X_subset and y_subset as input  
to the validation curve function (instead of X_mush and  
y_mush) to reduce computation time.**
```

```
#
```

```
# The initialized unfitted classifier object we'll be  
using is a Support Vector Classifier with radial basis  
kernel.
```

```
So your first step is to create an `SVC` object with  
default parameters (i.e. `kernel='rbf', C=1`) and  
`random_state=0`.
```

```
Recall that the kernel width of the RBF kernel is  
controlled using the `gamma` parameter.
```

```
# With this classifier, and the dataset in X_subset,
y_subset, explore the effect of `gamma` on classifier
accuracy by using the `validation_curve` function to
find the training and test scores for 6 values of
`gamma` from `0.0001` to `10` (i.e. `np.logspace(-
4,1,6)`).
```

You can specify what scoring metric you want `validation_curve` to use by setting the "scoring" parameter. In this case, we want to use "accuracy" as the scoring metric.

```
#
```

```
# For each level of `gamma`, `validation_curve` will
fit 3 models on different subsets of the data,
returning two 6x3 (6 levels of gamma x 3 fits per
level) arrays of the scores for the training and test
sets.
```

```
#
```

```
# Find the mean score across the three models for each
level of `gamma` for both arrays, creating two arrays
of length 6, and return a tuple with the two arrays.
```

```
#
```

```
# e.g.
```

```
#
```

```
# if one of your array of scores is
```

```
#
```

```
#      array([[ 0.5,  0.4,  0.6],
#             [ 0.7,  0.8,  0.7],
#             [ 0.9,  0.8,  0.8],
#             [ 0.8,  0.7,  0.8],
#             [ 0.7,  0.6,  0.6],
#             [ 0.4,  0.6,  0.5]])
```

```

#
# it should then become
#
#      array([ 0.5,  0.73333333,  0.83333333,
# 0.76666667,  0.63333333, 0.5])
#

# *This function should return one tuple of numpy
arrays `(training_scores, test_scores)` where each
array in the tuple has shape `(6,)`.

# In[131]:

def answer_two():
    from sklearn.svm import SVC
    from sklearn.model_selection import
validation_curve

    svc = SVC(kernel='rbf', C=1, random_state=0)
    gamma = np.logspace(-4,1,6)
    train_scores, test_scores = validation_curve(svc,
X_subset, y_subset,

                                                param_name='gamma',
                                                param_range=gamma,
                                                scoring='accuracy')

    scores = (train_scores.mean(axis=1),
test_scores.mean(axis=1))

    return scores # Your answer here

# In[137]:

```



```
print(answer_two())
```

```
for index, num in enumerate(np.logspace(-4,1,6)):  
    print(num)
```